

## Assignment 5 – PageRank Algorithm

**PageRank:** PageRank is a web graph ranking algorithm that helps Internet users sort hits by their importance.

PageRank calculates a numerical value for each element of a hyperlinked set of web pages, which reflects the probability that a random surfer will access that page.

PageRank is a recursive join plus update of rank values plus aggregation of ranks.

**Formula for PageRank:**

$$PR(p_i) = \frac{1-\alpha}{N} + \alpha \sum_{p_j \text{ links to } p_i} \frac{PR(p_j)}{L(p_j)}$$

**MapReduce:**

**Steps:-**

**Mapper:** The master nodes takes the input, divide it into smaller sub-problems and distributes them to worker nodes.

```
public class FinMapper extends Mapper<LongWritable, Text, DoubleWritable, Text> {

    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException,
        IllegalArgumentException {
        String line = value.toString(); // Converts Line to a String
        /*
         * TODO output key:-rank, value: node
         * See IterMapper for hints on parsing the output of IterReducer.
         */

        String[] sections = line.split("\t");
        String[] parts = sections[0].split(";");
        Double rank = Double.parseDouble(parts[1]);
        context.write(new DoubleWritable(-rank), new Text(parts[0]));
    }
}
```

**Reducer:** Reducer reduces a set of intermediate values which share a key to a smaller set of values.

```
public class FinReducer extends Reducer<DoubleWritable, Text, Text, Text> {

    public void reduce(DoubleWritable key, Iterable<Text> values, Context context)
        throws IOException,
        InterruptedException {
        /*
         * TODO: For each value, emit: key:value, value:-rank
         */
    }
}
```

```

        Iterator<Text> iterator = values.iterator();
        String node;
        while(iterator.hasNext()) {
            node = iterator.next().toString();
            context.write(new Text(node), new Text(String.valueOf(0 - key.get())));
        }
    }
}

```

#### InitMapper:

```

public class InitMapper extends Mapper<LongWritable, Text, Text, Text> {

    public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException,
            IllegalArgumentException {
        String line = value.toString(); // Converts Line to a String

        /*
         * TODO: Just echo the input, since it is already in adjacency list
        format.
         */

        String[] pair = line.split(":");
        if(pair != null && pair.length == 2) {
            context.write(new Text(pair[0].trim()), new Text(pair[1]));
        }
    }
}

```

#### InitReducer:

```

public class InitReducer extends Reducer<Text, Text, Text, Text> {

    public void reduce(Text key, Iterable<Text> values, Context context) throws
        IOException, InterruptedException {
        /*
         * TODO: Output key: node+rank, value: adjacency list
         */
        int defaultRank = 1;
        Iterator<Text> v = values.iterator();
        while(v.hasNext()) {
            context.write(new Text(key + "+" + defaultRank), v.next());
        }
    }
}

```

#### Iterator Mapper:

```

public class IterMapper extends Mapper<LongWritable, Text, Text, Text> {

    public void map(LongWritable key, Text value, Context context) throws
        IOException, InterruptedException,
            IllegalArgumentException {
        String line = value.toString();
        String[] sections = line.split("\t");
    }
}

```

```

        if (sections.length > 2) {
            throw new IOException("Incorrect data format");
        }
        if (sections.length != 2) {
            return;
        }
        /*
         * TODO: emit key: adj vertex, value: computed weight.
         *
         * Remember to also emit the input adjacency list for this node!
         * Put a marker on the string value to indicate it is an adjacency list.
         */
        String[] noderank = sections[0].split("\\+");
        String n = String.valueOf(noderank[0]);

        double rank = Double.valueOf(noderank[1]);
        String alist = sections[1].toString().trim();
        String[] anodes = alist.split(" ");

        int N = anodes.length;
        double weightOfPage = (double)1/N * rank;

        for(String adjacentnode : anodes) {
            context.write(new Text(adjacentnode), new
Text(String.valueOf(weightOfPage)));
        }
        context.write(new Text(n), new Text(PageRankDriver.MARKER +
sections[1]));
    }
}

```

#### Iterator Reducer:

```

public class IterReducer extends Reducer<Text, Text, Text, Text> {

    public void reduce(Text key, Iterable<Text> values, Context context) throws
IOException, InterruptedException {
        //double d = PageRankDriver.DECAY; // Decay factor
        /*
         * TODO: emit key:node+rank, value: adjacency list
         * Use PageRank algorithm to compute rank from weights contributed by
incoming edges.
         * Remember that one of the values will be marked as the adjacency list
for the node.
         */

        int defualtrank = 1;
        Iterator<Text> v = values.iterator();
        while(v.hasNext()) {
            context.write(new Text(key + "+" + defualtrank), v.next());
        }
    }
}

```

**Shuffle:** Input to the Reducer is the sorted output of the mappers. In this phase the framework fetches the relevant partition of the output of all the mappers, via http.

**Sort:** The framework groups Reducer inputs by keys in this stage.

**Partitioner:** Partitioner controls the partitioning of the keys of the intermediate map-outputs.

**Composite Function:**

```
public static void composite(String input, String output, String interim1,
                             String interim2, String diff, int reducers) throws Exception {
    /*
     * TODO
     */
    System.out.println("Keyur Doshi(10405923)");

    int counter = 0;
    init(input, interim1, reducers); // Initializes data
    counter++;
    double difference = 100000000; // sets ridiculously large default
difference
    int i = 0; // counter variables
    while (difference >= 30) // continues operation till difference is of
specified value
    {
        if (i % 2 == 0) {
            iter(interim1, interim2, reducers);
            /*
             * or every even iteration, interim1 is the input and
interim2 the output
             */
        } else {
            iter(interim2, interim1, reducers);
            // For odd iterations it is swapped
        }
        if (i % 3 == 0) // for every 3 iterations calculates the
difference
        {
            counter++;
            diff(interim1, interim2, "diffout", reducers);
            /*
             * as difference is an absolute value, order of input1,
input2 directories doesn't matter
             */
            difference = readDiffResult("diffout"); // Reads
difference from "diffout"
            System.out.println("Difference updates to:" + difference);
            deleteDirectory("diffout"); // deletes diffout
        }

        if (i % 2 == 0) // on even iterations, delete input directory
interim1
        {
            deleteDirectory(interim1);
        }
    }
}
```

```

        if (i % 2 == 1) // on odd, delete interim2
        {
            deleteDirectory(interim2);
        }
        counter++;
        i++; // increment i
    }

    if (i % 2 == 1) // As i increments at the last step, for odd i, interim2
                    // is the input directory
    {
        deleteDirectory(interim1); // deletes other directory
        counter++;
        finish(interim2, output, reducers);
        summarizeResult(output);
    } else
    {
        deleteDirectory(interim2); // Deletes other directory
        counter++;
        finish(interim1, output, reducers);
        summarizeResult(output);
    }
    System.out.println();
    System.out.println(counter);
}

```

### Here are the steps to setup standalone Hadoop cluster.

Install Java 8 :

```

sudo add-apt-repository ppa:webupd8team/java
Run Update : sudo apt-get update
sudo apt-get install oracle-java8-installer

```

Check java version:

```

java -version

```

Creating a Hadoop User for accessing HDFS and Mapreduce: (To avoid security issues, it is recommend to setup new Hadoop user group)

```

sudo addgroup hadoop
sudo adduser --ingroup hadoop huser
Enter password : - Create password to access huser
optional ignore
press Y

```

Minimize terminal and open new terminal:

```

Go to root directly: sudo su root
cd (change directory)

```

Giving correctly and securely root privileges to hadoop user "hhdhuser" by editing the /etc/sudoers file:

```

sudo gedit /etc/sudoers

```

Under user privileges  
huser ALL=(ALL:ALL) ALL

Install as well as configure SSH  
sudo apt-get install openssh-server

Hadoop uses SSH to access its nodes. For single node setup of Hadoop, we therefore need to configure SSH access to localhost.

Login with huser  
sudo su huser  
cd  
ssh-keygen -t rsa -P "" -- Generate SSH key for huser account  
ssh-keygen -t rsa -P ""

[If asked for a filename just leave it blank and press the enter key to continue]

Add the newly created key to the list of authorized keys so that hadoop can use SSH without prompting for a password.  
cat \$HOME/.ssh/id\_rsa.pub >> \$HOME/.ssh/authorized\_keys

Since hadoop doesn't work on IPv6 we should disable it. for that you need to update /etc/sysctl.conf  
# disable ipv6  
net.ipv6.conf.all.disable\_ipv6 = 1  
net.ipv6.conf.default.disable\_ipv6 = 1  
net.ipv6.conf.lo.disable\_ipv6 = 1

Command to reboot machine:  
sudo reboot

Download Hadoop hadoop-2.7.0.tar.gz  
Copy to desktop and extract

Terminal:  
sudo su huser  
cd

since the huser has root privileges, we can move the hadoop installation to the /usr/local/hadoop directory without any problem.  
sudo mv '/home/keyur/Desktop/hadoop-2.7.3' /usr/local/hadoop

Assign ownership of this folder to Hadoop user:  
sudo chown huser:hadoop -R /usr/local/hadoop

Create hadoop temp directories for Namenode and Datanode  
sudo mkdir -p /usr/local/hadoop\_tmp/hdfs/namenode  
sudo mkdir -p /usr/local/hadoop\_tmp/hdfs/datanode

Again assign ownership of this hadoop temp folder to hadoop user  
sudo chown huser:hadoop -R /usr/local/hadoop\_tmp

```
sudo gedit .bashrc
```

Add the following lines at the end of the file bashrc file

```
# -- HADOOP ENVIRONMENT VARIABLES START -- #
export JAVA_HOME=/usr/lib/jvm/java-8-oracle
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
export PATH=$PATH:$HADOOP_HOME/sbin
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib"
export PATH=$PATH:/usr/local/hadoop/bin/
# -- HADOOP ENVIRONMENT VARIABLES END -- #
```

Now we going to configure hadoop-env.sh - This file contains some environment variable settings used by Hadoop such as where log files are stored, the maximum amount of heap used etc.

```
cd /usr/local/hadoop/etc/hadoop
sudo gedit hadoop-env.sh
```

```
Comment export JAVA_HOME = ${JAVA_HOME}
```

Add or change the vairable JAVA\_HOME which specifies the path to the Java Installation used by hadoop

```
export JAVA_HOME='/usr/lib/jvm/java-8-oracle'
```

Configuration of core-site.xml file: core-site.xml contains configuration information that overrides the default values for core Hadoop properties.

```
sudo gedit core-site.xml
```

```
## Name node runs on localhost on port 9000
<property>
<name>fs.default.name</name>
<value>hdfs://localhost:9000</value>
</property>
```

Configuration of hdfs-site.xml file: The hdfs-site.xml file enables you to overwrite a number of default values that control the HDFS

```
sudo gedit hdfs-site.xml
```

```
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>file:/usr/local/hadoop_tmp/hdfs/namenode</value>
</property>
<property>
```

```
<name>dfs.datanode.data.dir</name>
<value>file:/usr/local/hadoop_tmp/hdfs/datanode</value>
</property>
```

\*dfs.replication - The default replication factor to use for each block of a file.

\*dfs.namenode.name.dir - The namenode stores the metadata of the file system in the local file system in the specified directory.

\*dfs.datanode.name.dir - The datanode stores the actual blocks of file data in the local file system in the specified directory.

Configuration of yarn-site.xml: yarn-site.xml file enables you to overwrite a number of default values controlling the YARN components.  
sudo gedit yarn-site.xml

```
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
```

\*yarn.nodemanager.aux-services - The name of an auxiliary service being added to the Node Manager.

\*yarn.nodemanager.aux-services.mapreduce\_shuffle.class - Implements the mapreduce\_shuffle service

Configuration of yarn-site.xml: yarn-site.xml file enables you to overwrite a number of default values controlling the YARN components.  
cp /usr/local/hadoop/etc/hadoop/mapred-site.xml.template  
/usr/local/hadoop/etc/hadoop/mapred-site.xml

Configuration of mapred-site.xml: mapred-site.xml file enables you to overwrite a number of default values controlling the Mapreduce job execution components.  
sudo gedit mapred-site.xml

## This runtime framework that is used to execute MapReduce jobs.

```
<property>
<name>mapreduce.framework.name</name>
<value>yarn </value>
</property>
```

```
cd
clear
```

Prepare to start the hadoop cluster by formatting the namenode, which needs to be done only once.

```
hdfs namenode -format or
/usr/local/hadoop/bin/hadoop namenode -format
```

```
cd /usr/local/hadoop
```



```
Start hdfs daemons/services  
start-dfs.sh or  
sbin/start-dfs.sh
```

### **Format the New Hadoop Filesystem**

```
hduser@ubuntu:~$ hadoop namenode --format
```

### **Starting Hadoop**

```
hduser@ubuntu:~$ Start-all.sh
```

### **Stopping Hadoop**

```
hduser@ubuntu:~$ Stop-all.sh
```

### **Checking if it's really up and running:**

```
hduser@ubuntu:~$ jps
```